

F21_GPU_Proj01 Report

Jiale Shi

I take the following public resources as my reference to do this job.
Without the help of these public resources, I could not finish this homework.

1. <https://stackoverflow.com/questions/39336574/box-filter-with-cuda-c>
2. https://www.nvidia.com/content/nvision2008/tech_presentations/Game_Developer_Track/NVISION08-Image_Processing_and_Video_with_CUDA.pdf

The project is a box linear filter for image processing.

https://en.wikipedia.org/wiki/Box_blur

Sunflower_3.jpg is the output after box linear filter.

CUDA kernel launch with (174,122) blocks of (16,16) threads in my code.

I find that the matr_multiply.cu cannot include pointer_2d_matrix.h.
Therefore, I cannot call Matrix in matr_multiply.cu.
In order to solve the problem, I redefine arrays in the main.cpp and pass array to the cuda_kernel_wrapper().

```
float *data_Red, *data_Blue, *data_Green;  
data_Red = (float *)malloc(Width*Height*sizeof(float));  
data_Blue = (float *)malloc(Width*Height*sizeof(float));  
data_Green = (float *)malloc(Width*Height*sizeof(float));
```

I pass the array individually to cuda_kernel_wrapper() to do box linear filter.

```
// this eventually calls cuda kernel  
cout << "Process Red:" << endl;  
cuda_kernel_wrapper(data_Red, Width, Height);
```

```
cout << "Process Blue:" << endl;
cuda_kernel_wrapper(data_Blue, Width, Height);
cout << "Process Green:" << endl;
cuda_kernel_wrapper(data_Green, Width, Height);
```

cuda_kernel_wrapper() uses the d_filter kernel.

```
__global__ void d_filter(float *g_idata, float *g_odata, unsigned int width,
unsigned int height)
```

Inside the d_filter, we use shared memory.

Performance Test:

I cut the original Sunflower.jpg to create smaller images.

Sunflower.jpg

Image width: 2434 pixels

Image height: 1697 pixels

Process Red:

CUDA kernel launch with 174,122 blocks of 16,16 threads

Time to run the kernel: **0.325792 ms**

Process Blue:

CUDA kernel launch with 174,122 blocks of 16,16 threads

Time to run the kernel: **0.276480 ms**

Process Green:

CUDA kernel launch with 174,122 blocks of 16,16 threads

Time to run the kernel: **0.281760 ms**

Sunflower_M.jpg

Image width: 1930 pixels

Image height: 1396 pixels



Process Red:

CUDA kernel launch with 138,100 blocks of 16,16 threads

Time to run the kernel: **0.231648 ms**

Process Blue:

CUDA kernel launch with 138,100 blocks of 16,16 threads

Time to run the kernel: **0.185824 ms**

Process Green:

CUDA kernel launch with 138,100 blocks of 16,16 threads

Time to run the kernel: **0.184064 ms**

Sunflower_S.jpg

Image width: 1537 pixels

Image height: 1063 pixels



Process Red:

CUDA kernel launch with 110,76 blocks of 16,16 threads

Time to run the kernel: **0.156512 ms**

Process Blue:

CUDA kernel launch with 110,76 blocks of 16,16 threads

Time to run the kernel: **0.113984 ms**

Process Green:

CUDA kernel launch with 110,76 blocks of 16,16 threads

Time to run the kernel: **0.106848 ms**

Sunflower_XS.jpg
Image width: 644 pixels
Image height: 601 pixels



Process Red:
CUDA kernel launch with 47,43 blocks of 16,16 threads
Time to run the kernel: **0.063040 ms**

Process Blue:
CUDA kernel launch with 47,43 blocks of 16,16 threads
Time to run the kernel: **0.031680 ms**

Process Green:
CUDA kernel launch with 47,43 blocks of 16,16 threads
Time to run the kernel: **0.028992 ms**

Conclusion:
Smaller size images need smaller blocksPerGrid (fewer blocks) and the time to run the kernel would be shorter.

